Quinn Anderson

quinn.anderson@email.com | +1-415-555-0117 | San Francisco, CA

PROFESSIONAL SUMMARY

MLOps Engineer with 6 years building ML platforms and infrastructure. Expert in model serving, monitoring, and scalable ML systems architecture.

SKILLS

Python | Kubernetes | Docker | MLflow | Kubeflow | Ray | Seldon | BentoML | Prometheus | Grafana | Apache Airflow | Terraform | AWS | Feature Stores | Model Monitoring

PROFESSIONAL EXPERIENCE

Senior MLOps Engineer - MLPlatform Co. (2020 - Present)

- Built internal ML platform serving 100+ models with automated A/B testing.
- Reduced model deployment time from 2 weeks to 2 hours through automation.
- Implemented model drift detection reducing incidents by 80%.

ML Infrastructure Engineer - AIScale (2018 - 2020)

- Developed feature store supporting 500+ features for 20 ML models.
- Optimized model inference latency by 70% using ONNX and TensorRT.

EDUCATION

M.S. Computer Science | Georgia Tech (2018)

B.S. Computer Science | UIUC (2016)